



DSA 104 AI and ML in Chemistry

Session 3: Introduction to ML – Definitions and Supervised Models

Dr. Johannes Schögenhumer (johannes.schoergenhumer@chem.uzh.ch)



```
import tensorflow as tf  
  
model.add(Dense(64, activation='relu'))  
  
optimizer = tf.keras.optimizers.Adam(),  
  
model.compile(loss='categorical_crossentropy',  
  
model.fit(X_train, y_train, epochs=10)
```

AI – ML – DL: Definition and Distinction?

AI (Artificial Intelligence):

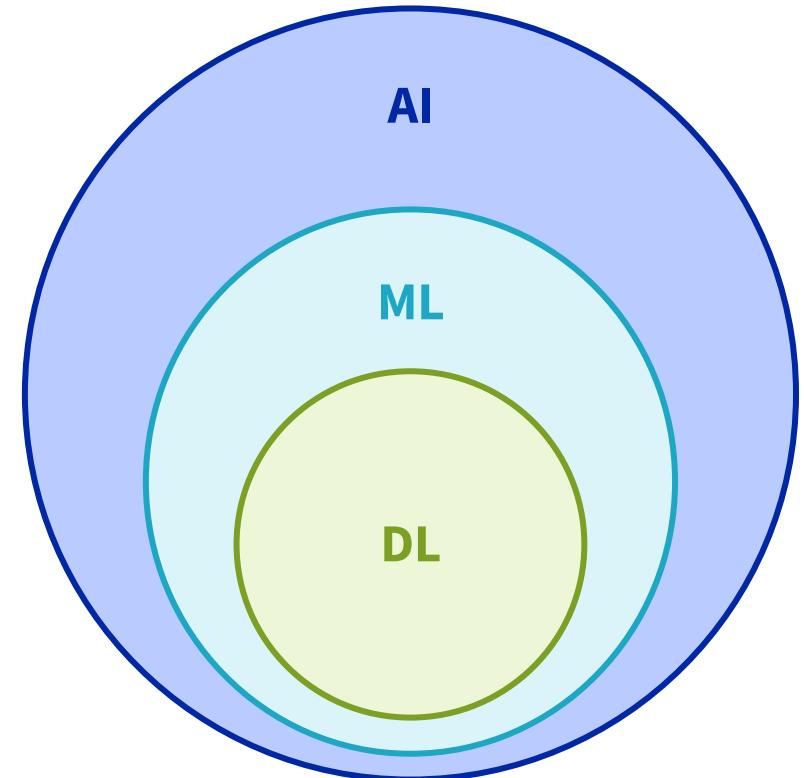
- A broad “umbrella term” for systems behaving intelligently
- Not all AI involves data!
- Historically includes e.g. symbolic reasoning, rule-based systems, logic systems (e.g. Classic chess engines using search, Logic-based planning systems)

ML (Machine Learning):

- Learn patterns from data to improve performance on a task
- Replaces explicit programming with learning from examples, **rules inferred from data**
- Includes e.g. Regressors, Random Forest, GradientBoost, DeepLearning!

DL (Deep learning):

- Learn hierarchical representations using neural networks with multiple layers.
- Instead of hand-designing features, DL learns features automatically.
- Subset of ML, extremely powerful for complex data (high data demand)



Exercise: Classification quiz – AI, ML or DL?

- AI** 1. A thermostat
- AI** 2. Siri or Alexa recognizing your voice command
- DL** 3. ChatGPT
- AI** 4. A chatbot responding to customer service queries using rules
- ML** 5. Netflix recommender
- DL** 6. Generating realistic images using GANs (Generative Adversarial Networks)
- AI** 7. Expert system for medical diagnosis
- DL** 8. CNN for image recognition
- ML** 9. Detecting spam emails using a classifier trained on labeled emails
- AI** 10. A chess computer
- DL** 11. Self-driving car detecting pedestrians

Types of ML

Supervised Learning:

- Labelled data
- Direct feedback
- Predict outcome

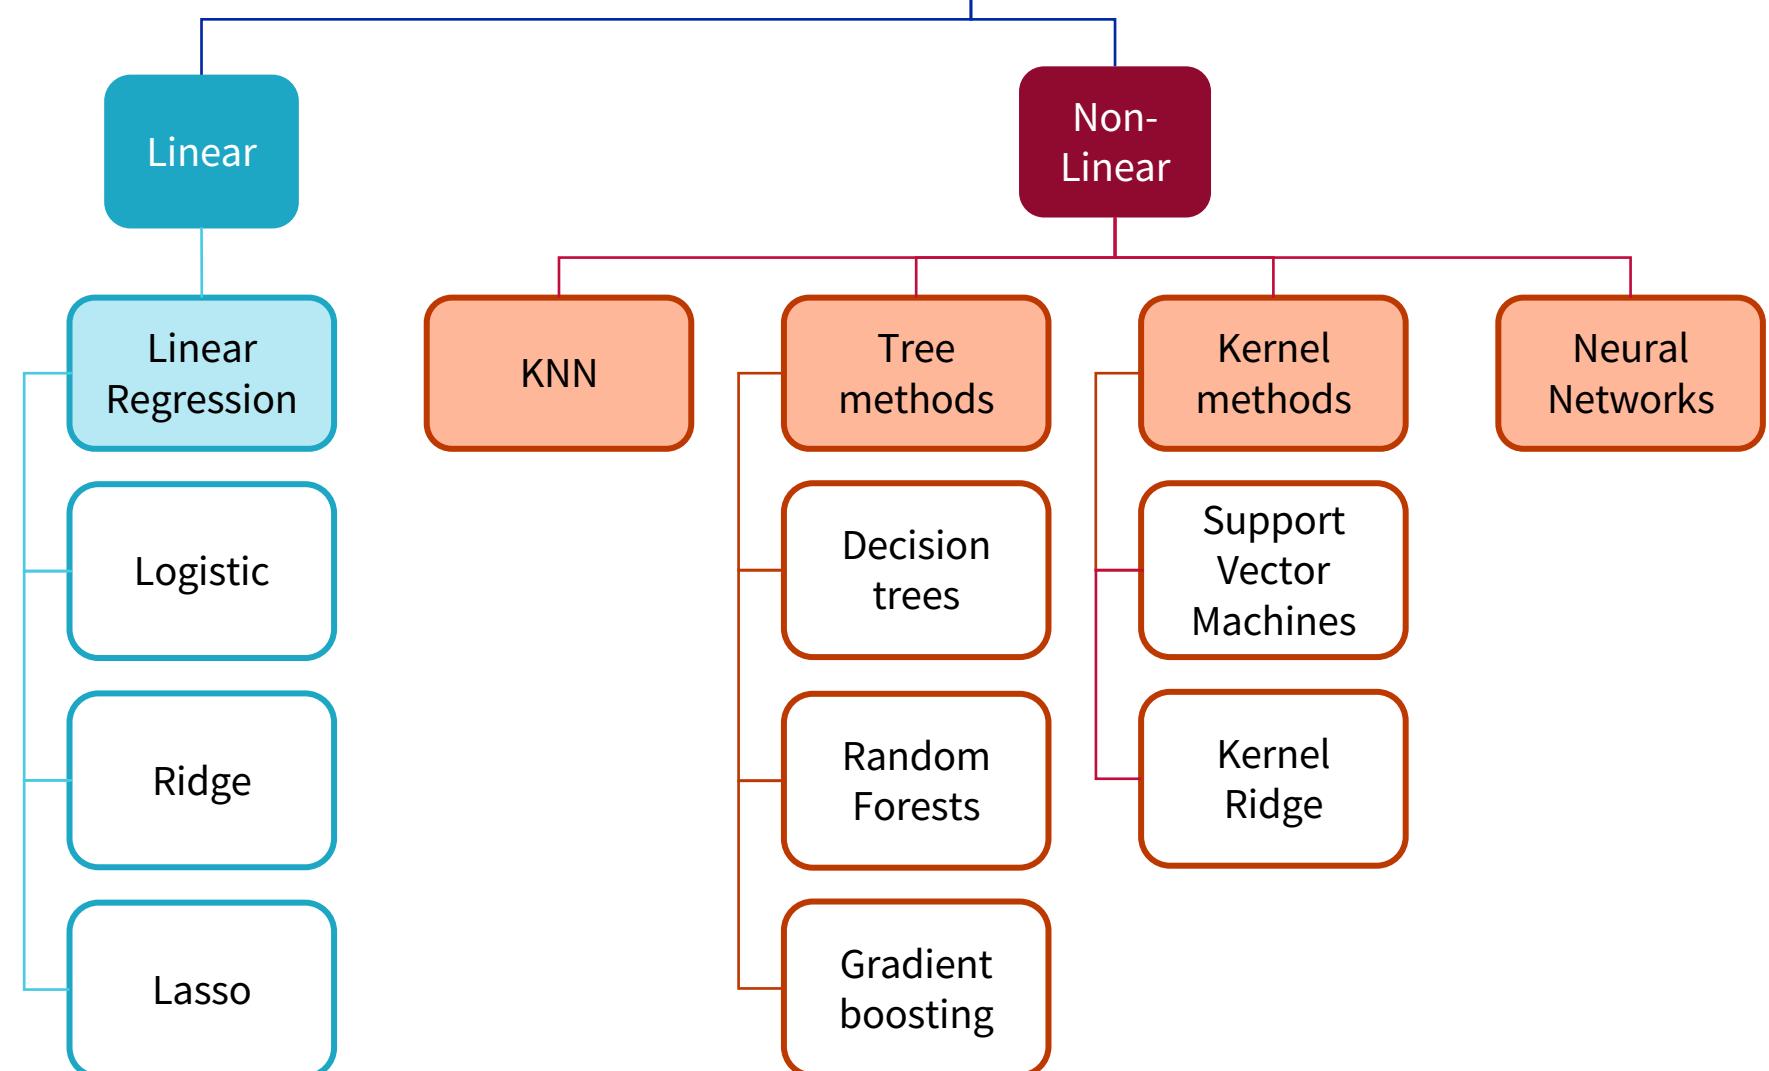
Unsupervised Learning:

- No labels/targets
- No feedback
- Find hidden structure in data

Reinforcement learning:

- Decision process
- Reward system
- Learn series of actions

Overview of Supervised ML algorithms



Supervised ML

Learning from **input values** and corresponding **target values**:

- E.g. image + object type, DNA sequence + phenotype, ...

Typical usage: **predictive modelling**

- **Train** model on data set with input + target values.
- Use trained model to **predict** target values for other (new) inputs where the targets are not known yet.
- Classic predictions:
 - **Classification**: target value is class **label**
 - **Regression**: target value is **numerical** value

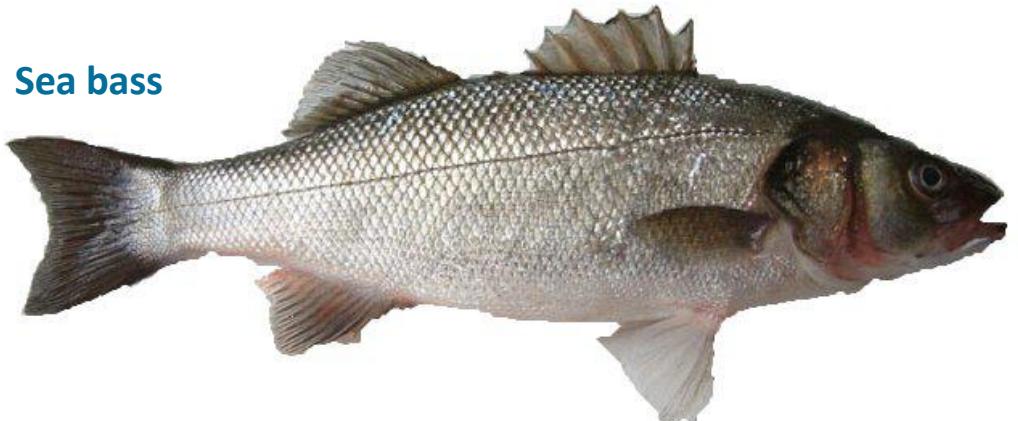
Example:

- Automated system to sort fish in a fish-packing company: salmons must be distinguished from sea bass optically.

Salmon



Sea bass



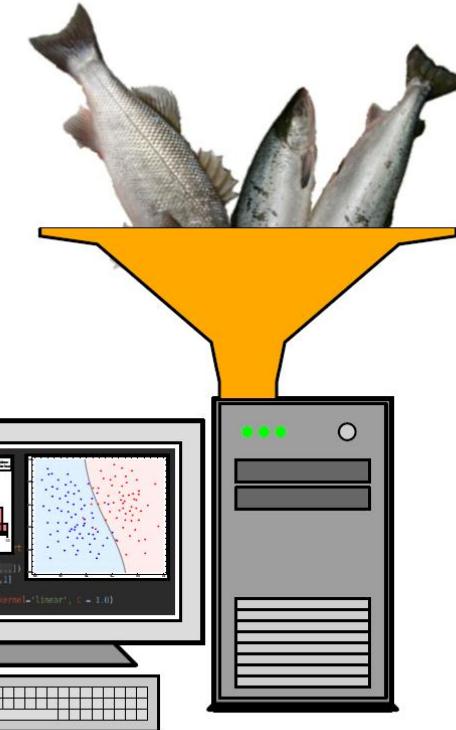
Example from R. O. Duda, P. E. Hart, and D. G. Stork. Pattern Classification. 2nd edition. John Wiley & Sons, 2001.
ISBN 0-471-05669-3.

Terminology: model

- **Model:** parameterized function/method with specific parameter values (e.g. a trained neural network)
- **Model class:** the class of models in which we search for the model (e.g. neural networks, SVMs, ...)
- **Parameters:** what is adjusted during training (e.g. network weights)
- **Hyperparameters:** settings controlling model complexity or the training procedure (e.g. network learning rate)
- **Model selection/training:** process of finding a model (optimal parameters) from the model class

Example:

- Automated system to sort fish in a fish-packing company: salmons must be distinguished from sea bass optically.



Terminology: data

We can represent an object by a **vector x** of feature values (=feature vector) of **length d** and **label y** :

$$x = (x_1, \dots, x_d) \quad y$$

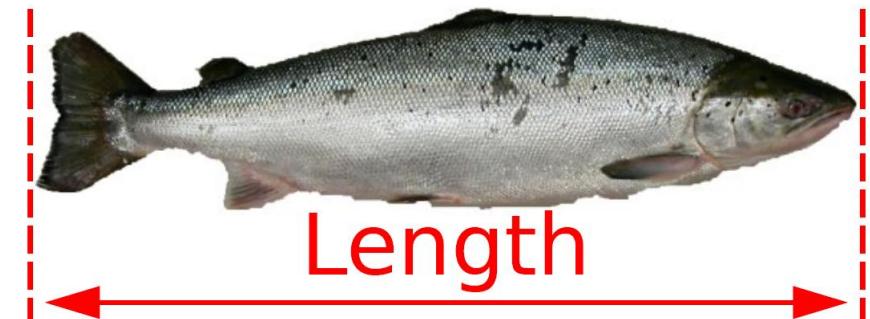
- A fish is represented as feature vector with **two values length and brightness** (i.e. $d = 2$) and **one label** ($y = \text{"salmon"}$ or $y = \text{"sea bass"}$)
- An object described by one feature vector and one label is referred to as **sample**: (x, y) .
- n samples each with feature vectors x_1, \dots, x_n of length d and label y_1, \dots, y_n , thus, **matrix of feature vectors X** and the labels in a corresponding **labels vector y** :

$$X = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} x_{11} & \cdots & x_{1d} \\ \vdots & \ddots & \vdots \\ x_{n1} & \cdots & x_{nd} \end{bmatrix} \quad y = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}$$

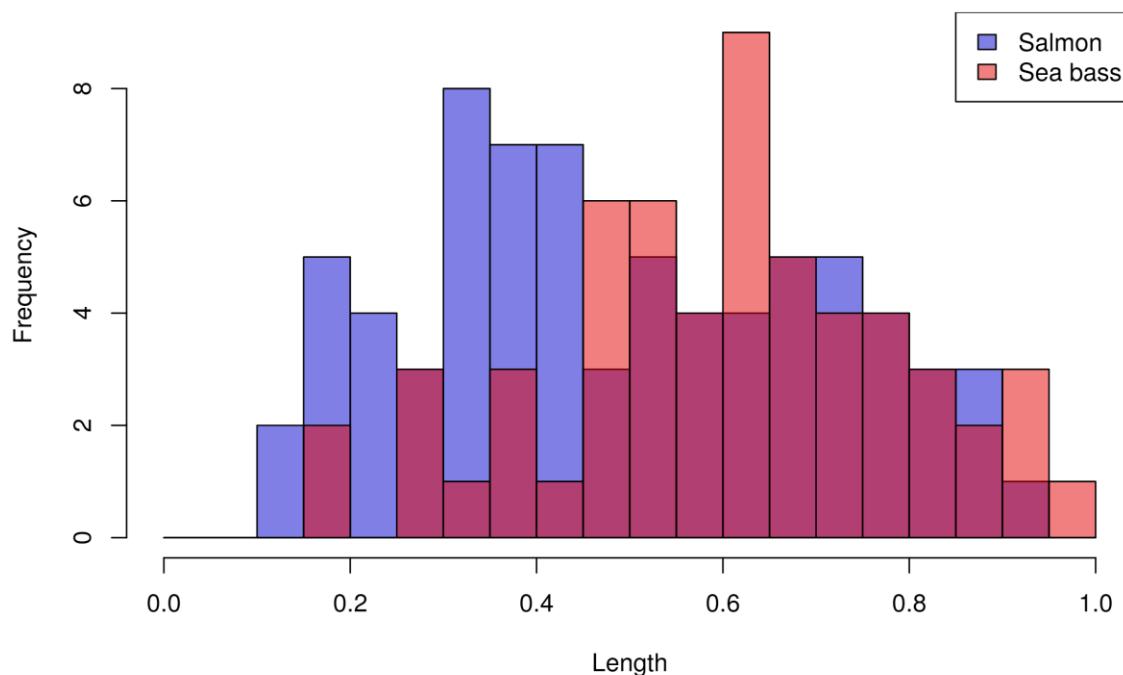
- Our labeled data is thus described by: (X, y) .

Example:

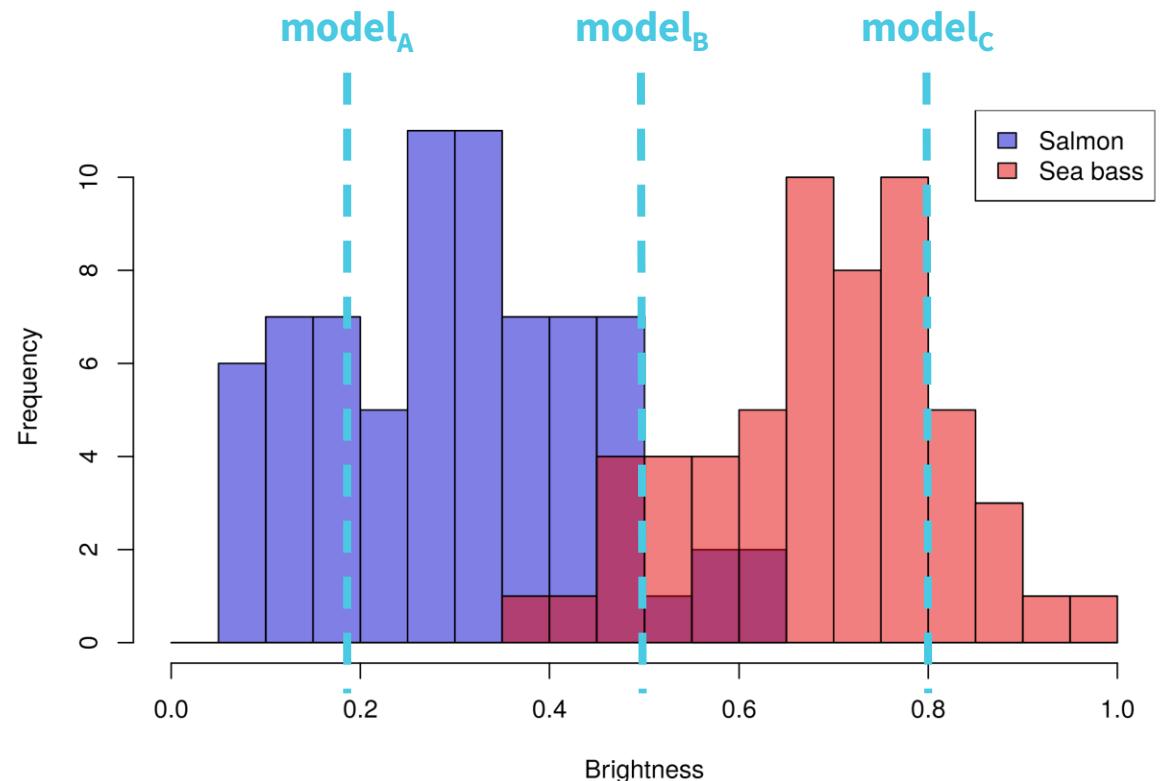
Salmon



Fish example: Look at the data



- Brightness looks more useful for fish classification.
- 3 different models based on brightness threshold:
 - model_A: brightness < 0.18 → Salmon
 - model_B: brightness < 0.5 → Salmon
 - model_C: brightness < 0.8 → Salmon



How do we get the “best” model?

- How does our model perform on our data?
 - **Loss function**
- How will it perform on (unseen) future data?
 - **Generalisation error / risk**

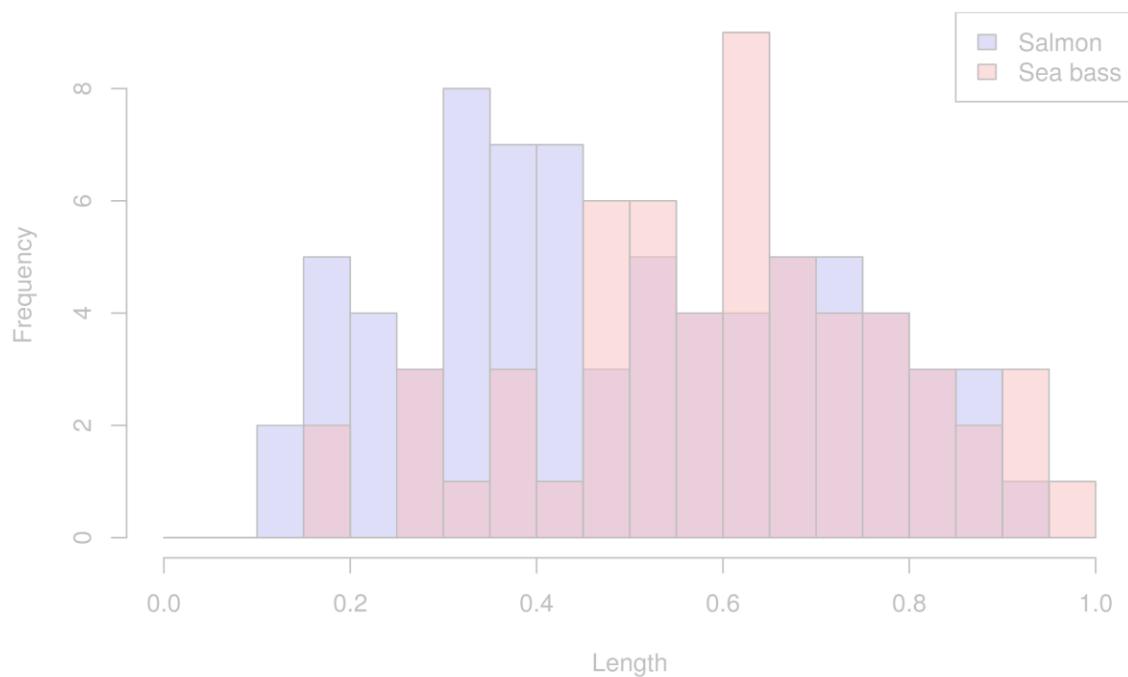
Loss functions

- Assume we have a model \mathbf{g} with parameter \mathbf{w} :
 - $\mathbf{g}(x; \mathbf{w})$ maps an input vector x to an output value \hat{y}
 - \hat{y} needs to be as close as possible to the true target value y
- We can use a **loss function L** to measure how close the prediction is to the true target for a sample with (x, y) :
 - $L(y, \mathbf{g}(x; \mathbf{w})) = L(y, \hat{y})$
 - Quantifies the error: Lower loss = better model performance
 - Adjust the parameters/weights \mathbf{w} to reduce loss

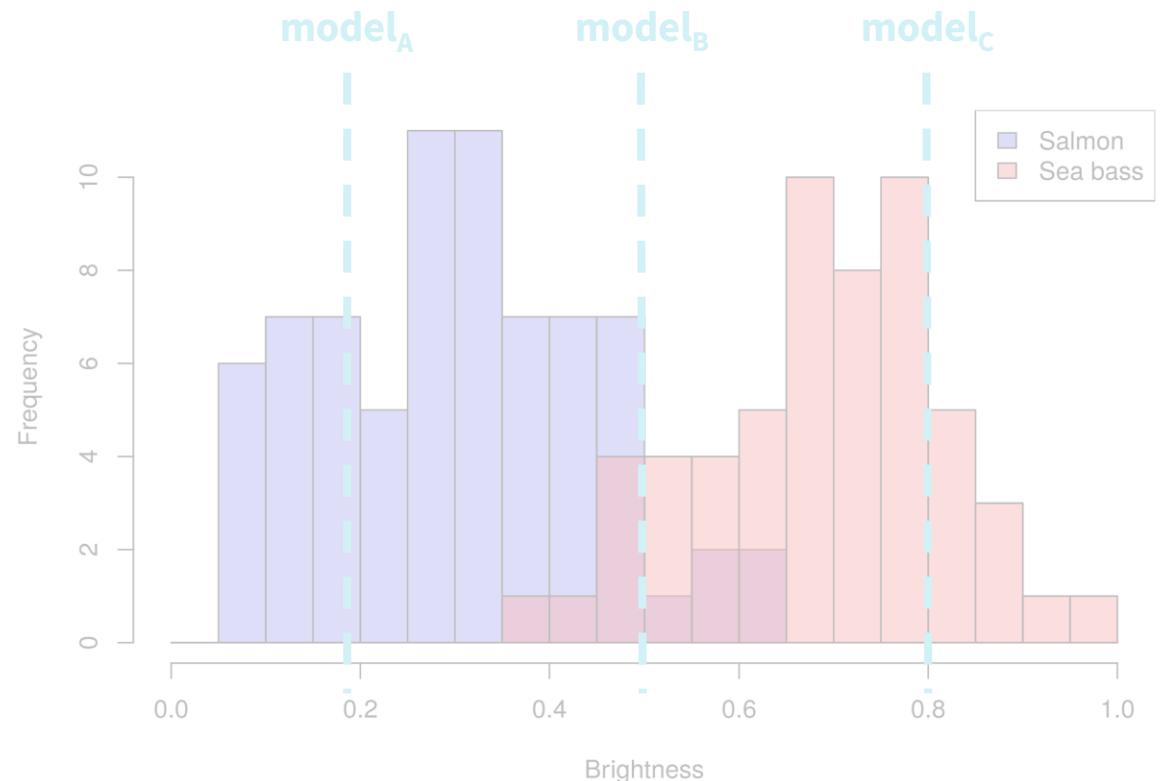
Loss functions: Examples

- Zero-one loss: $L_{\text{ZO}}(y, g(x; w)) = \begin{cases} 0 & y = g(x; w) \\ 1 & y \neq g(x; w) \end{cases}$ Classification errors
- MSE: $L_{\text{MSE}}(y, g(x; w)) = \frac{1}{n} \sum_{i=1}^n (y_i - g(x_i; w))^2$ Quite common, heavily punishes outliers (as compared to MAE), not for robustness
- Many other loss functions available
- **Different loss functions are needed for different data, tasks and model classes, e.g.**
 - Regression: Mean Squared Error (MSE), Mean Absolute Error (MAE)
 - Classification: Cross-Entropy Loss (log loss), Hinge Loss (SVMs)
 - Ranking/Boosting: Exponential loss (AdaBoost), Logistic loss (Gradient Boosting)

Fish example: Look at the data



- Brightness looks more useful for fish classification.
- 3 different models based on brightness threshold:
 - model_A: brightness < 0.18 → Salmon
 - model_B: brightness < 0.5 → Salmon
 - model_C: brightness < 0.8 → Salmon



How do we get the “best” model?

- How does our model perform on our data?
 - **Loss function** ✓
- How will it perform on (unseen) future data?
 - **Generalisation error / risk**

Generalisation error / risk

$$\hat{y} = g(x; w)$$

The **generalisation error** or **risk** is the expected loss on **future / unseen data** for a given model $g(x; w)$

$$R(g(x; w)) = \iint_{Xy} L(y, g(x; w)) \cdot p(x, y) \cdot dy dx$$

$R(g(x; w))$ denotes the **expected loss** for input x , and $p(x, y)$ is the joint probability distribution for x and y .

- In practice, joint probability distribution $p(x, y)$ unknown:
- **Estimate** the generalisation error: **Empirical Risk Minimisation (ERM)**

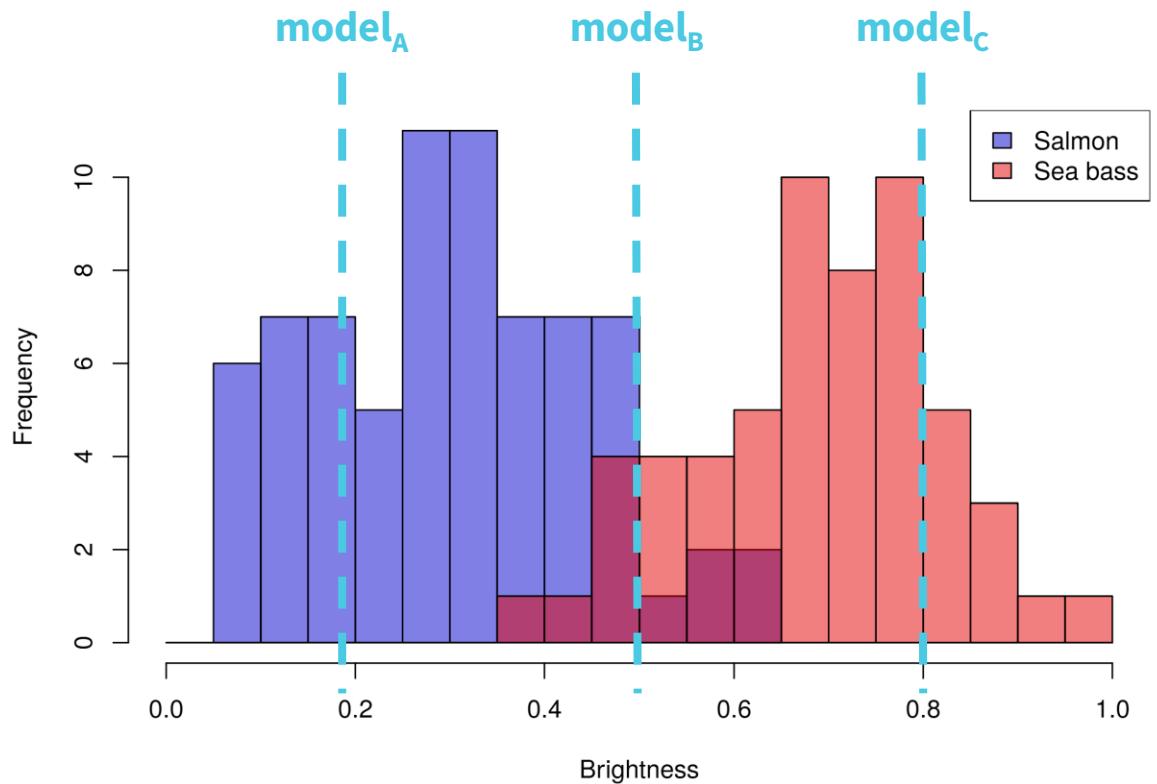
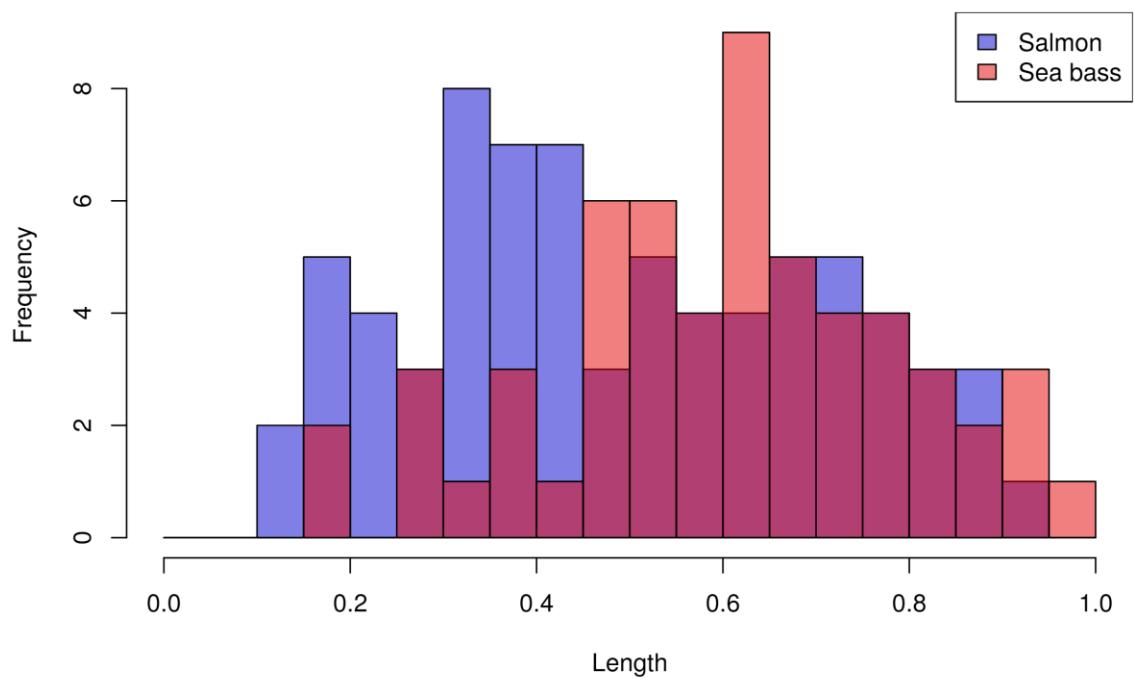
- Use data set (X, y) instead of $p(x, y)$:

$$R_E(g(x; w), (X, y)) = \frac{1}{n} \sum_{i=1}^n L(y_i, g(x_i; w))$$

- Law of large numbers: for $n \rightarrow \infty$: $R_E(g(x; w)) \rightarrow R(g(x; w))$

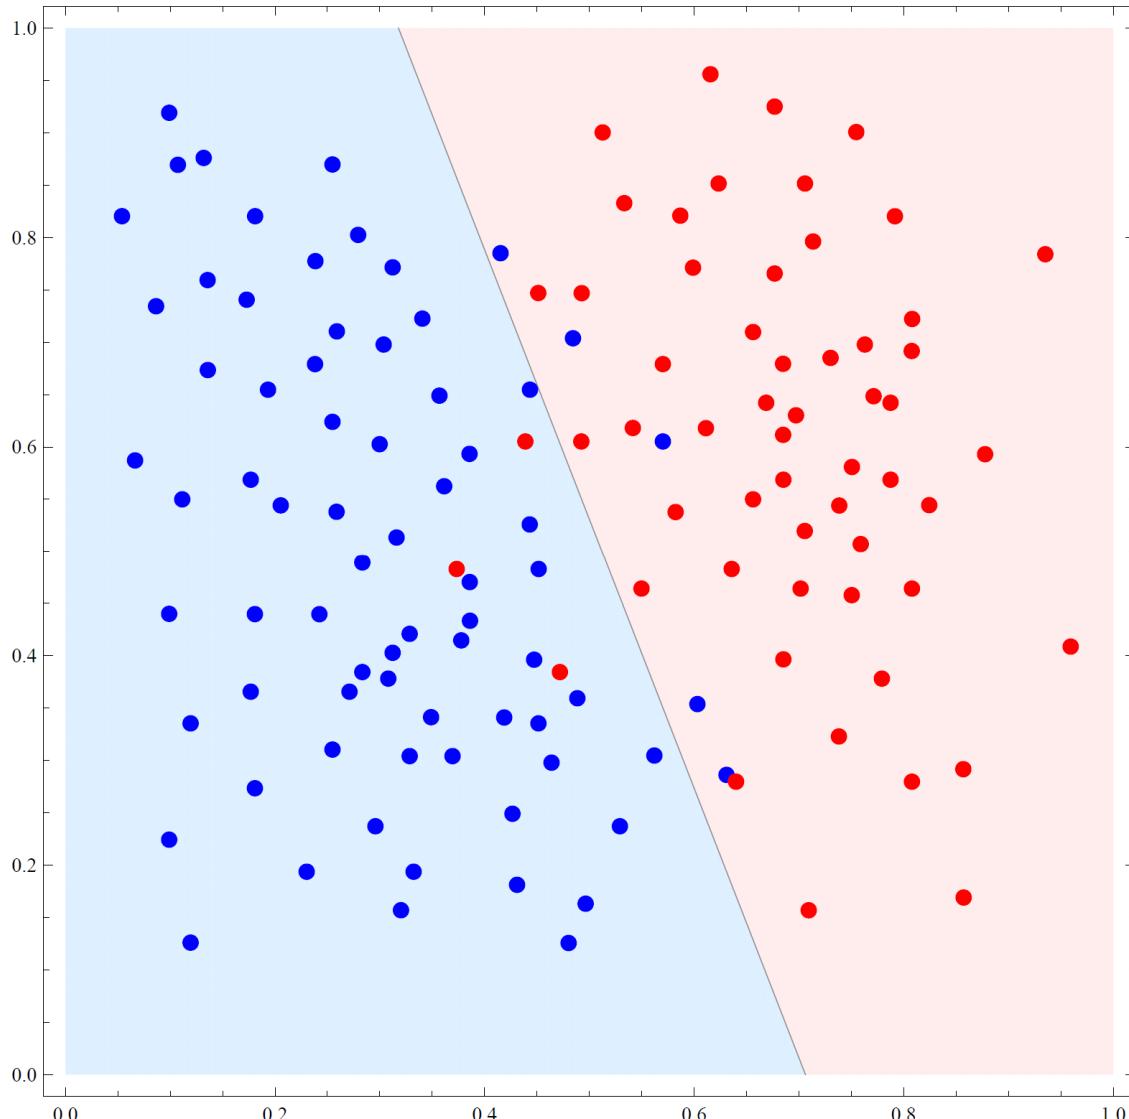
- Minimisation: $g_{\text{ERM}} = \arg \min_{g \in G} R_E(g)$

Fish example



ERM optimises model only based on training data.
Individual features (especially length) do not separate classes super well
Combine our features and use a different model class

Fish example – linear separation:

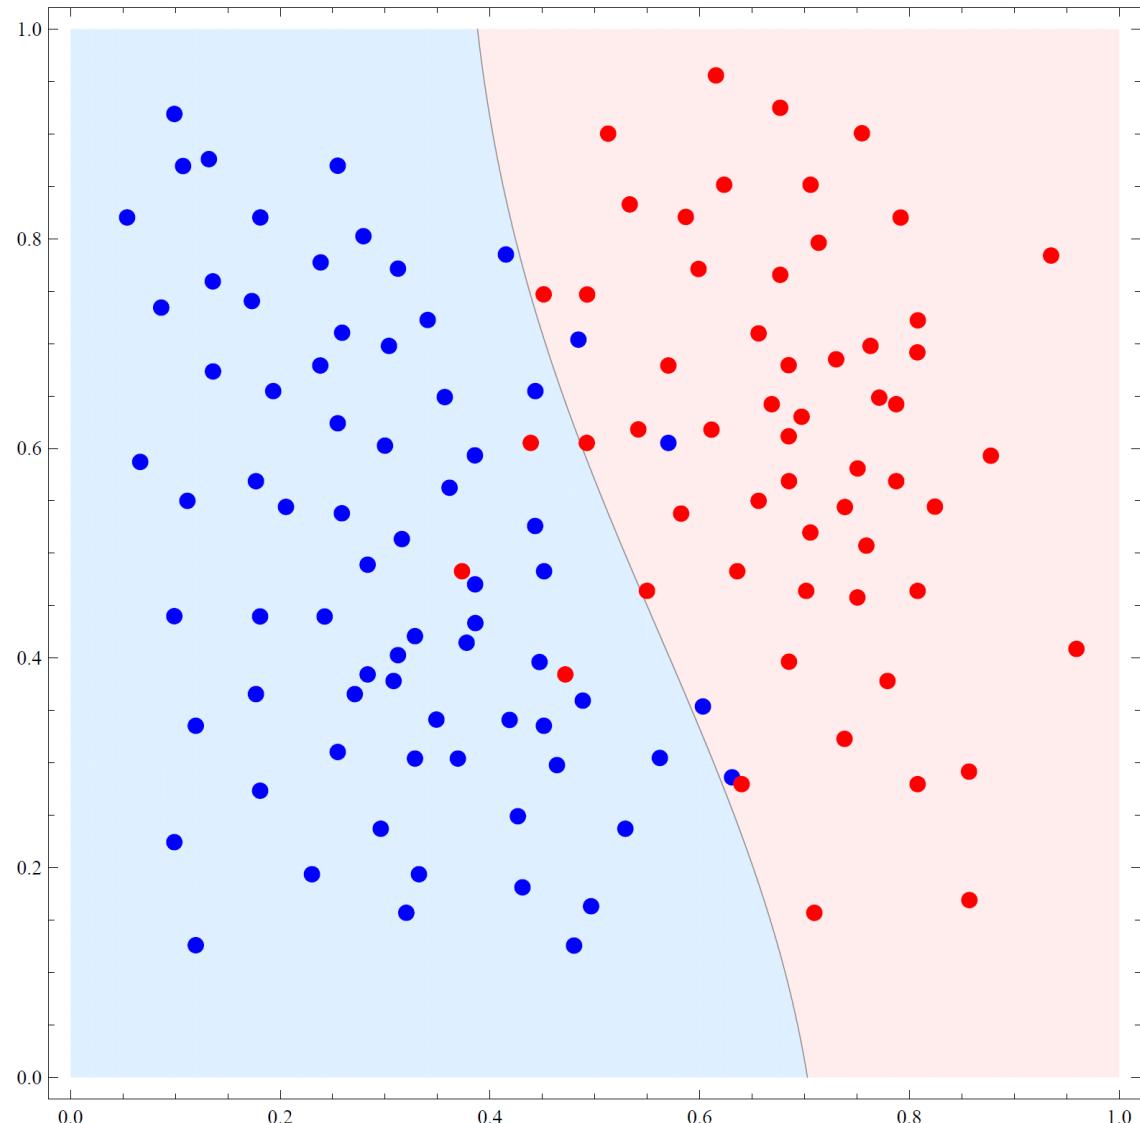


Adapted from A. Schörgenhuber, *Hands-on AI I*, Lecture materials, 2023, JKU Linz.

University of Zurich | DSA 104 AI and ML in Chemistry

23/02/2026 | 15

Fish example – mildly non-linear separation:

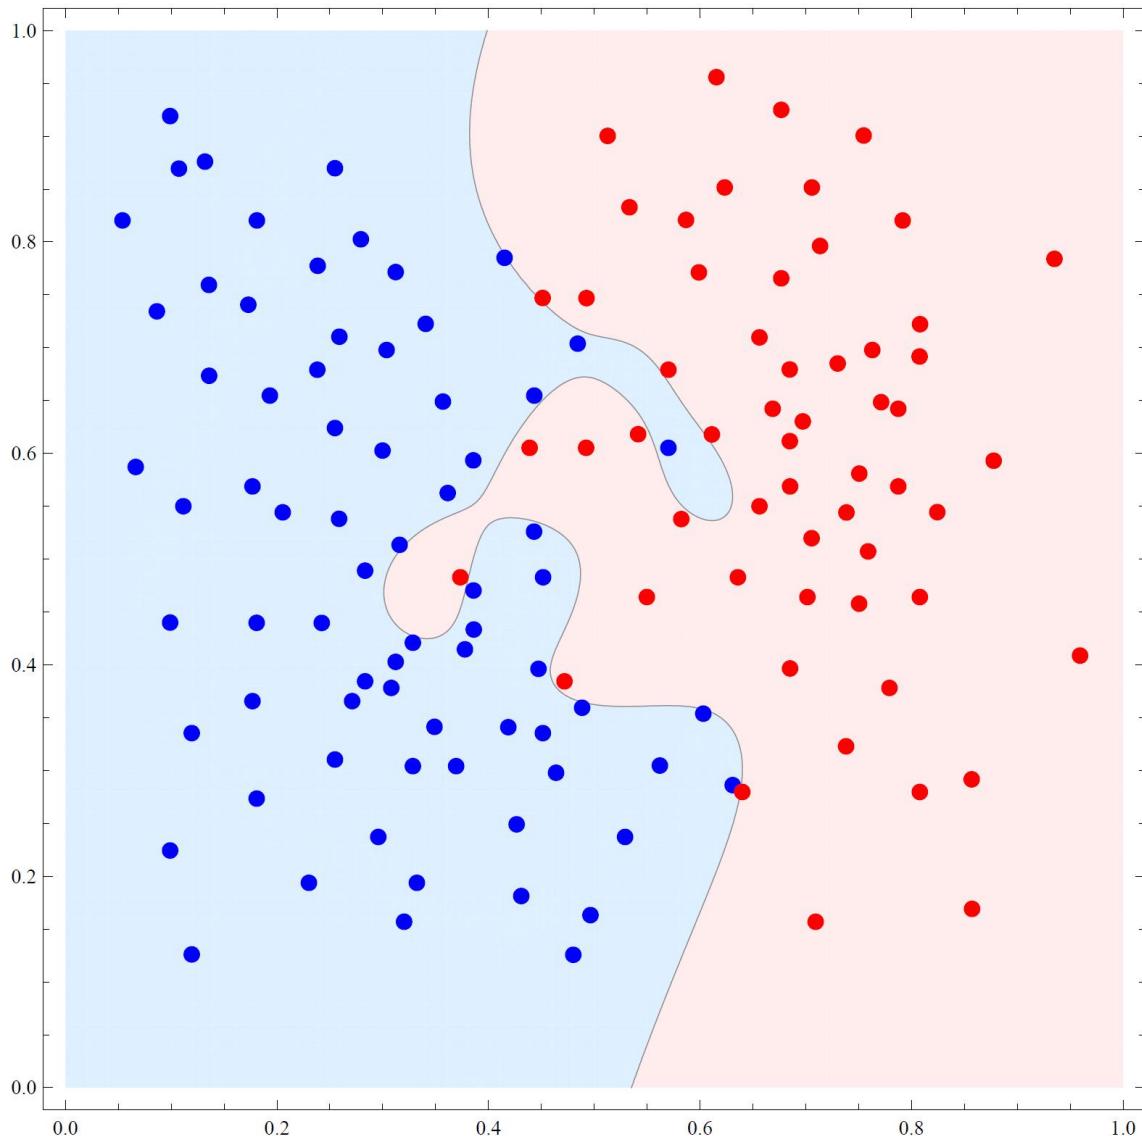


Adapted from A. Schörgenhofer, *Hands-on AI I*, Lecture materials, 2023, JKU Linz.

University of Zurich | DSA 104 AI and ML in Chemistry

23/02/2026 | 16

Fish example – highly non-linear separation:



Overfitting?

Better risk estimation

Test Set Method

- Assumption: Samples are **independently and identically distributed**
- Split data set of n samples into **two non-overlapping subsets**:
 - **Training set**: for ERM (i.e. to optimise parameters), typically 80% of n
 - **Test set**: use to estimate the risk (test data = approximation of future, unseen data), typically 20% of n
- RE on the test set identifies overfitting!

Better risk estimation

Validation Set

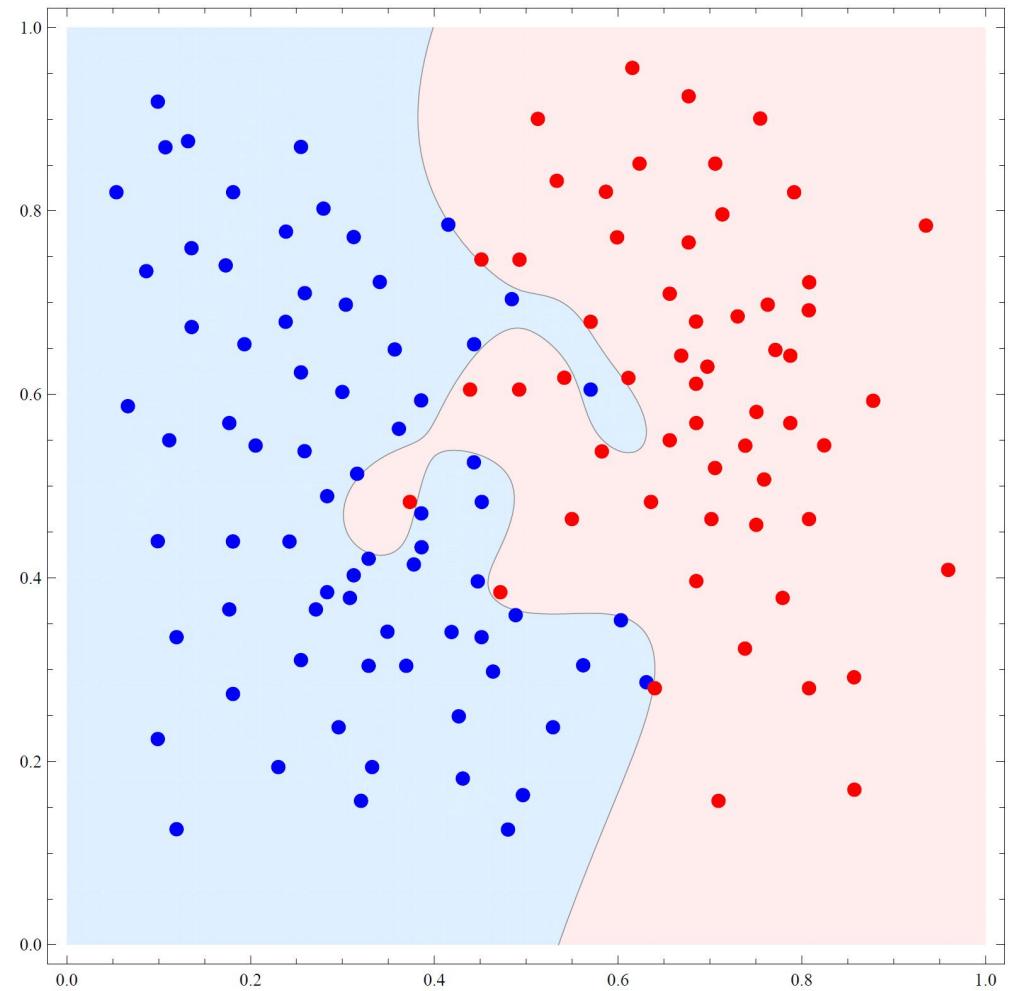
- Hyperparameters often changed, model will be retrained and evaluated (**hyperparameter tuning**)
- May lead to overfitting:
 - Hyperparameter tuning based on test set evaluation results → Test set indirectly used for training
 - Not non-overlapping subsets - estimation of the generalisation error compromised
- Solution: Create a third non-overlapping **validation set**.
- Use the validation set for hyperparameter tuning and the test set (once) for the final model evaluation.

More on better RE and Generalisation:

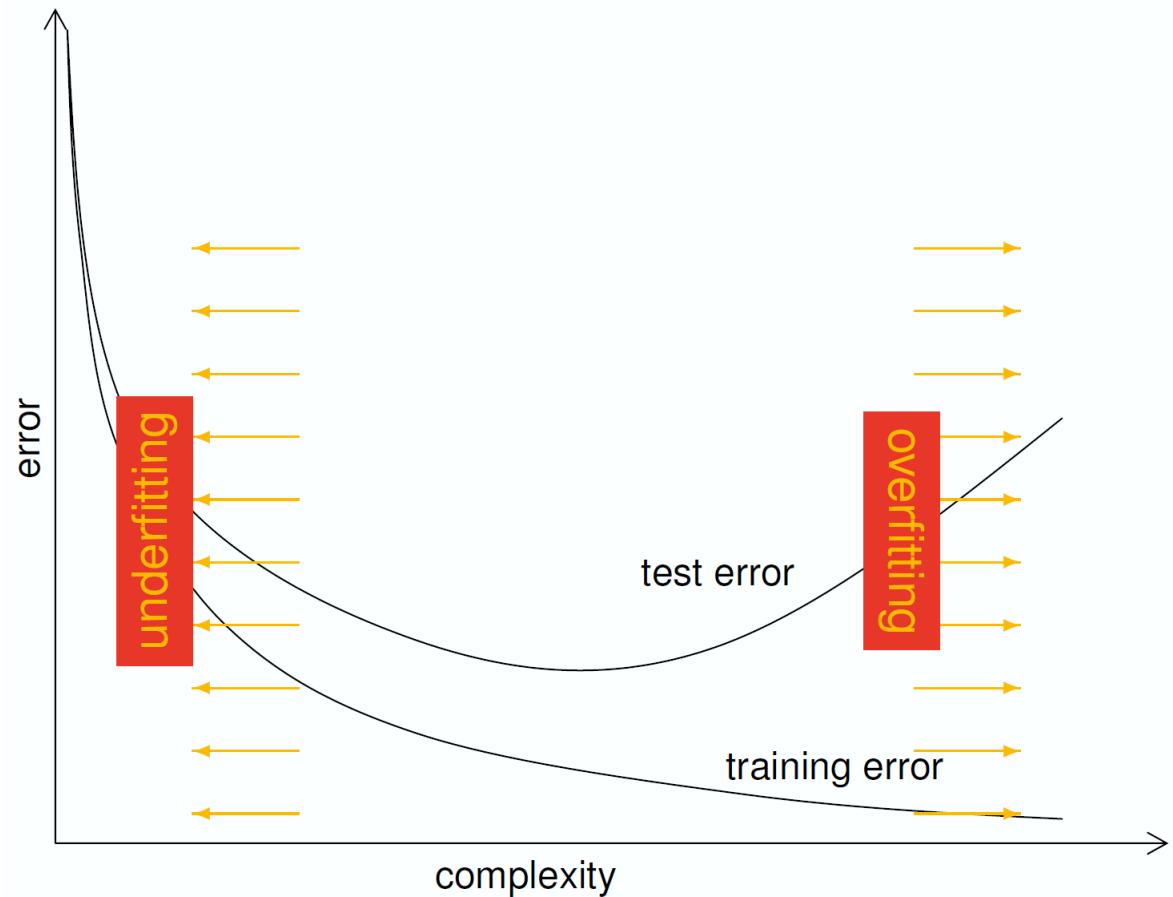
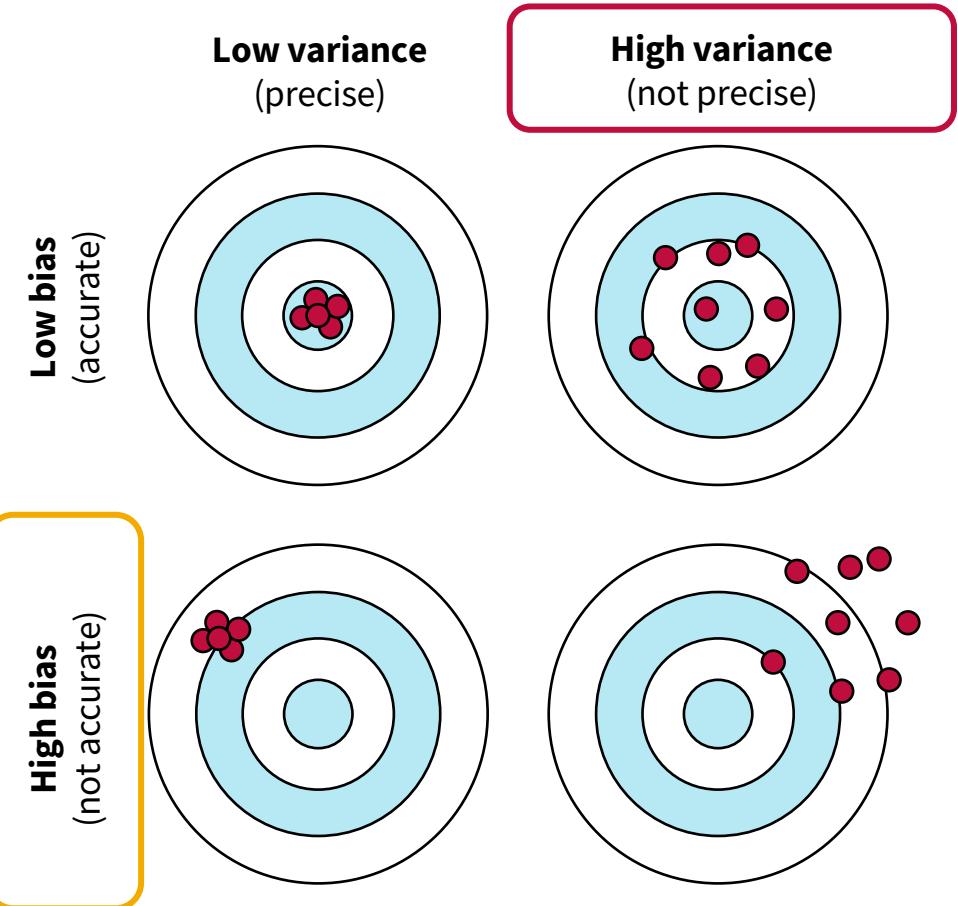
- Good generalisation requires combination of ERM with...
 - Regularisation (e.g. add penalties, limit parameters): $L' (y, \hat{y}) = L(y, \hat{y}) + \lambda \cdot \Omega(\hat{y})$...where $\Omega(\hat{y})$ is a **regularisation function**, e.g.
$$\Omega(\hat{y}) = ||\theta||_2^2 = w_1^2 + w_2^2 + \dots + w_j^2$$
 - Appropriate model capacity
 - **Enough data!!**
...with θ being a the parameter vector of j parameters w

Fish example – best fit

- Now, we can use **ERM** to optimise a model on our **training data set** (optionally, including a validation set)
- A held-out **test set** will allow us to get an estimate about the performance on future data (i.e. **generalisation**)
- If overfitting is detected, we can reduce the model complexity via hyperparameters
- But what if the model is **too simple**? → Underfitting!



Bias-variance tradeoff

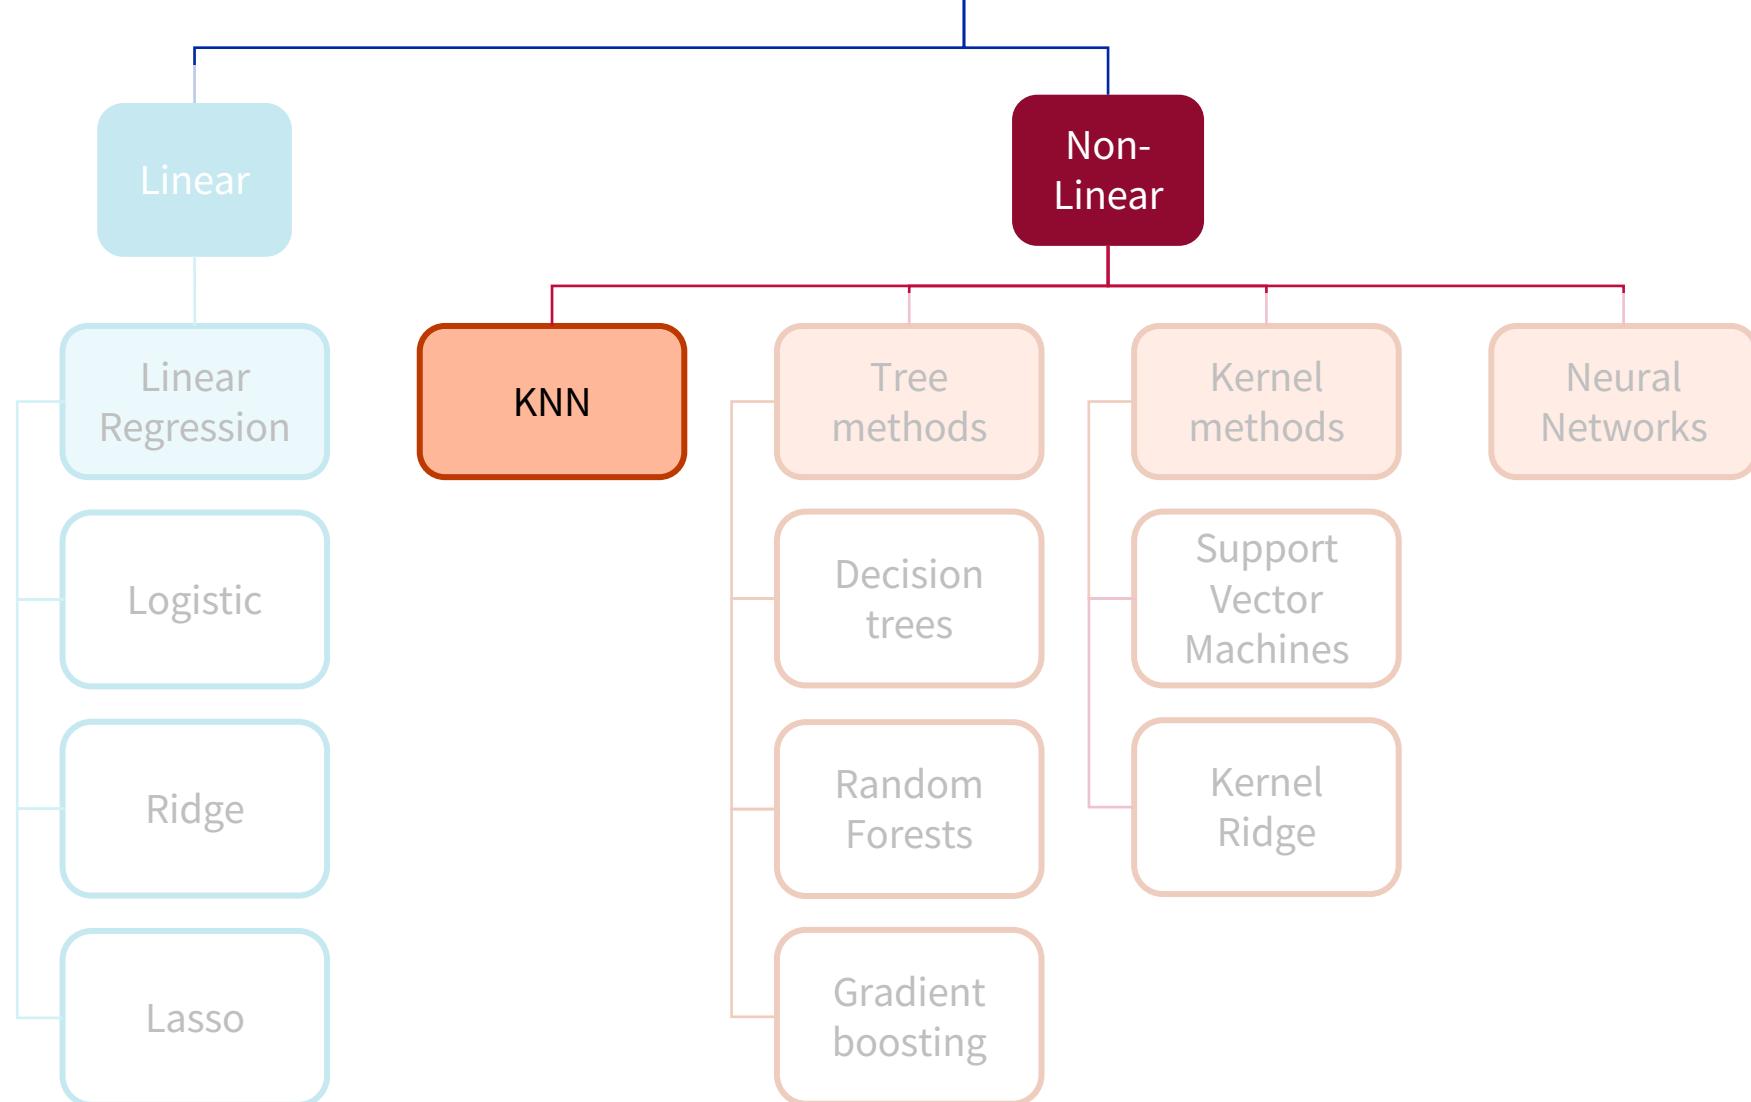


- **Underfitting** (high bias): The model is too coarse to fit training data and also too coarse to fit test data. The model complexity is too low.
- **Overfitting** (high variance): The model fits (too) well to training data but not to future/test data. The model complexity is too high.

Exercise: Over- or Underfitting? Or just right?

1. A linear regression model is used to predict house prices, but it consistently misses trends and has high error on both training and test sets. **Underfitting**
2. A decision tree model perfectly predicts the training data but performs worse on unseen data. **Overfitting**
3. A model achieves similar accuracy on both training and test data. **Well-fitted**
4. Using a polynomial regression of degree 5 on a small dataset, the model perfectly fits training points but fluctuates wildly on new data. **Overfitting**
5. A linear model predicts sales poorly because it cannot capture the underlying pattern. **Underfitting**
6. A neural network achieves 98% accuracy on the training set, 90% on the test set, but shows high variance between predictions on different validation batches **Overfitting**

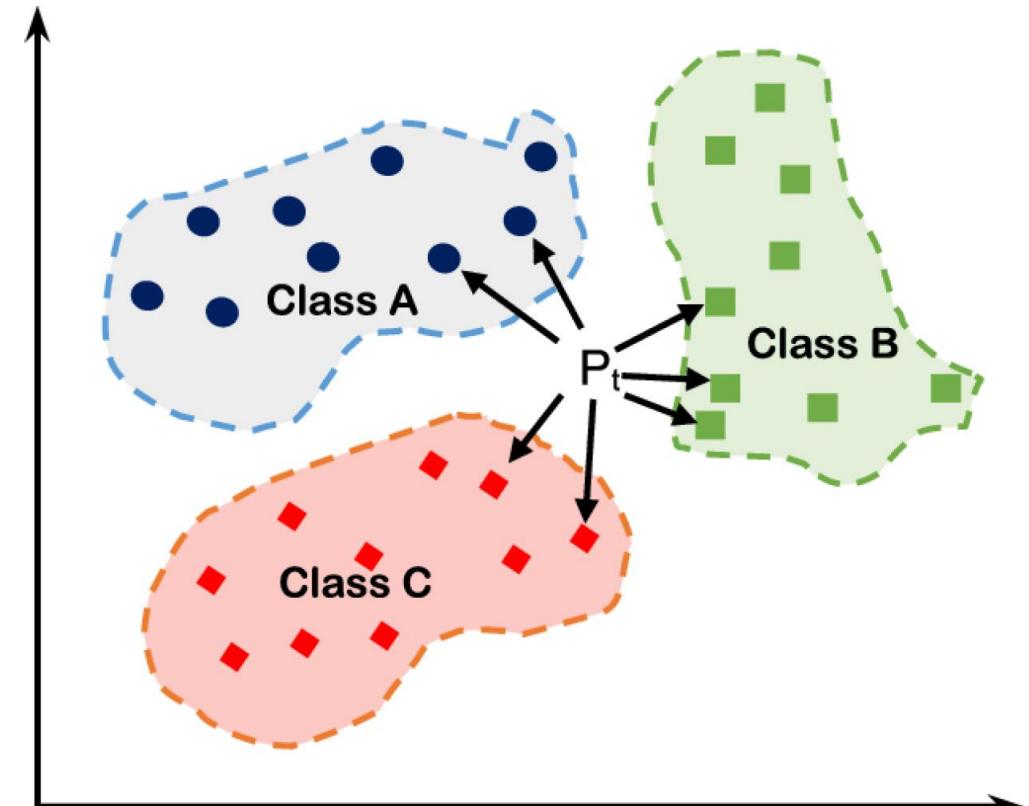
Overview of Supervised ML algorithms



k-Nearest Neighbors Classifier (KNN)

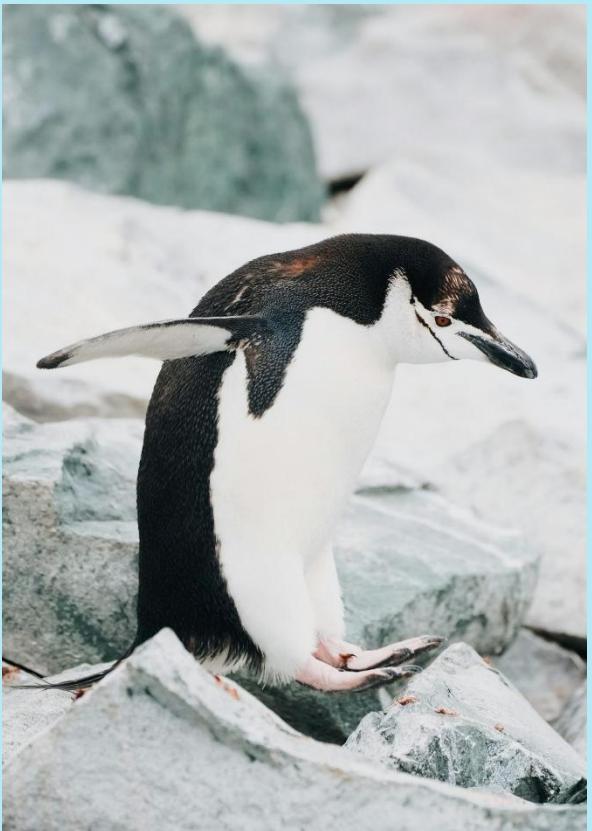
Assume we have a labeled data set (X, y) and a **distance measure** on the input space.

- The **k-nearest neighbors classifier** is defined as follows:
 - $g_{k\text{-NN}}(x; w)$ = class that occurs most often among k samples closest to x
- For $k = 1$: **nearest neighbor classifier**:
 - $g_{\text{NN}}(x; w)$ = class of the sample that is closest to x
- In case of ties: e.g. random class assignment or class with larger number of samples is assigned



Classification model KNN

- Explore the impact of different numbers for K (*KNN-penguins-decisionboundaries*)
- Use a KNN Classifier to predict different penguin species for unknown data! (*KNN-penguins*)



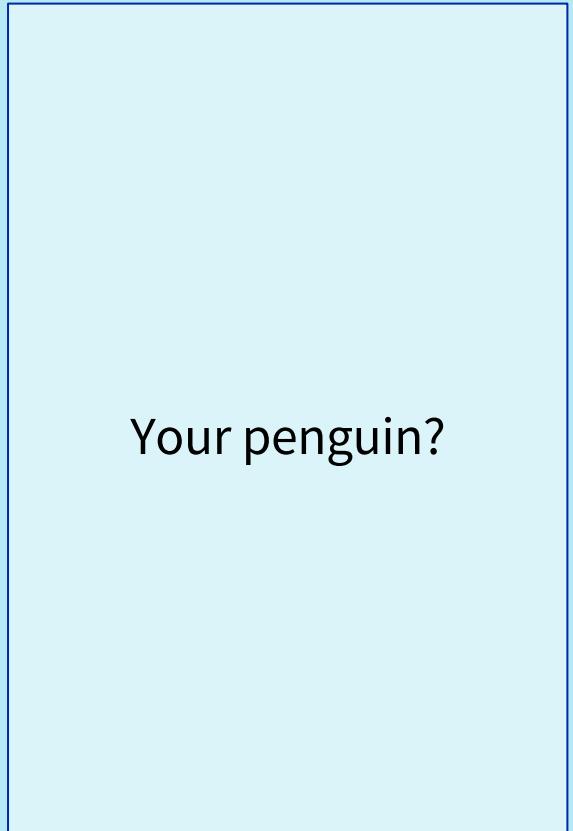
Chinstrap



Adélie



Gentoo



Your penguin?